C:\WINDOWS\SYSTEM32\cmd. | Windows PowerShell | Command Prompt - uv run ru

```
C:\Users\fady8>cd C:\Users\fady8\repos\UDACITY_MCP

C:\Users\fady8\repos\UDACITY_MCP>uv run run_queries.py
2026-07-07 19:57:35,887 - INFO - ✓ Server 'llm_inference' initialized
2026-07-07 19:57:36,314 - INFO - ✓ Server 'sqlite' initialized
2026-07-07 19:57:37,313 - INFO - ✓ Server 'filesystem' initialized
2026-07-07 19:57:37,319 - INFO - Attempting execution of tool 'write_query' (attempt 1/3)
2026-07-07 19:57:37,324 - INFO - ✓ Data extraction tables initialized

=== STARTING AUTOMATED QUERIES ===

Query 1: scrape these sites: {"cloudrift": "https://www.cloudrift.ai/inference", "deepinfra": "https://deepinfra.com/pri
cing", "fireworks": "https://fireworks.ai/pricing#serverless-pricing", "groq": "https://groq.com/pricing"}

2026-07-07 19:57:41,772 - INFO - HTTP Request: POST https://claude.vocareum.com/v1/messages "HTTP/1.1 200 OK"
I'll scrape these AI inference provider websites for you.
[Using tool: scrape_websites]
2026-07-07 19:57:41,781 - INFO - Attempting execution of tool 'scrape_websites' (attempt 1/3)
2026-07-07 19:57:46,475 - INFO - HTTP Request: POST https://claude.vocareum.com/v1/messages "HTTP/1.1 200 OK"
Great! I've successfully scraped all four websites:

1. **CloudRift** - https://www.cloudrift.ai/inference
2. **DeepInfra** - https://deepinfra.com/pricing
3. **Fireworks** - https://fireworks.ai/pricing#serverless-pricing
4. **Groq** - https://groq.com/pricing

All the content has been scraped and stored in both markdown and HTML formats. You can now use the `extract_scraped_info
` function with any of these provider names (cloudrift, deepinfra, fireworks, or groq) to retrieve and view the scraped
content.

Would you like me to extract and display the information from any of these scraped websites?2026-07-07 19:57:52,027 - IN
FO - HTTP Request: POST https://claude.vocareum.com/v1/messages "HTTP/1.1 200 OK"
2026-07-07 19:57:52,027 - INFO - Attempting execution of tool 'write_query' (attempt 1/3)
2026-07-07 19:57:52,039 - INFO - Attempting execution of tool 'write_query' (attempt 1/3)
2026-07-07 19:57:52,047 - INFO - Attempting execution of tool 'write_query' (attempt 1/3)
2026-07-07 19:57:52,057 - INFO - Attempting execution of tool 'write_query' (attempt 1/3)
2026-07-07 19:57:52,066 - INFO - Stored 4 pricing plans

Query 2: How much does cloudrift ai (https://www.cloudrift.ai/inference) charge for deepseek v3?

2026-07-07 19:57:54,772 - INFO - HTTP Request: POST https://claude.vocareum.com/v1/messages "HTTP/1.1 200 OK"
I'll extract the CloudRift pricing information to find the DeepSeek V3 pricing.
[Using tool: extract_scraped_info]
2026-07-07 19:57:54,772 - INFO - Attempting execution of tool 'extract_scraped_info' (attempt 1/3)
2026-07-07 19:57:58,557 - INFO - HTTP Request: POST https://claude.vocareum.com/v1/messages "HTTP/1.1 200 OK"
Based on the scraped information from CloudRift AI (https://www.cloudrift.ai/inference), here's the pricing for **DeepSe
ek V3**:

- **Input Price**: $0.27 per 1M tokens
- **Output Price**: $1.10 per 1M tokens

This is their serverless LLM inference pricing. They also offer DeepSeek R1 at $0.55 (input) / $2.19 (output) per 1M tok
ens.2026-07-07 19:58:02,990 - INFO - HTTP Request: POST https://claude.vocareum.com/v1/messages "HTTP/1.1 200 OK"
2026-07-07 19:58:02,991 - INFO - Attempting execution of tool 'write_query' (attempt 1/3)
2026-07-07 19:58:03,003 - INFO - Attempting execution of tool 'write_query' (attempt 1/3)
2026-07-07 19:58:03,012 - INFO - Stored 2 pricing plans

Query 3: How much does deepinfra (https://deepinfra.com/pricing) charge for deepseek v3

2026-07-07 19:58:05,457 - INFO - HTTP Request: POST https://claude.vocareum.com/v1/messages "HTTP/1.1 200 OK"
I'll extract the DeepInfra pricing information to find the DeepSeek V3 pricing.
[Using tool: extract_scraped_info]
2026-07-07 19:58:05,460 - INFO - Attempting execution of tool 'extract_scraped_info' (attempt 1/3)
2026-07-07 19:58:09,803 - INFO - HTTP Request: POST https://claude.vocareum.com/v1/messages "HTTP/1.1 200 OK"
Based on the scraped information from DeepInfra (https://deepinfra.com/pricing), here's the pricing for **DeepSeek V3**:

- **Input Price**: $0.32 per 1M tokens
- **Output Price**: $0.89 per 1M tokens
```

```
For a typical 50/50 input/output split (500K input, 500K output tokens):
- **CloudRift**: $0.135 + $0.55 = **$0.685**
- **DeepInfra**: $0.16 + $0.445 = **$0.605**
- **DeepInfra saves ~12%** in this scenario2026-07-07 19:58:51,053 - INFO - HTTP Request: POST https://claude.vocareum.c
om/v1/messages "HTTP/1.1 200 OK"
2026-07-07 19:58:51,056 - INFO - Stored 0 pricing plans

=== SHOWING STORED DATA ===
2026-07-07 19:58:51,056 - INFO - Attempting execution of tool 'read_query' (attempt 1/3)

==================================================
Stored Pricing Plans:
==================================================
- Company: DeepInfra, Plan: DeepSeek V3, Input: $0.32 USD/1M tokens, Output: $0.89 USD/1M tokens
==================================================
```

```
Query 4: Compare cloudrift ai and deepinfra's costs for deepseek v3


[Enforcing database read for comparison query]
2026-07-07 19:58:17,896 - INFO - HTTP Request: POST https://claude.vocareum.com/v1/messages "HTTP/1.1 200 OK"
I'll retrieve the pricing data from the database to compare CloudRift AI and DeepInfra's costs for DeepSeek V3.
[Using tool: read_query]
2026-07-07 19:58:17,896 - INFO - Attempting execution of tool 'read_query' (attempt 1/3)
2026-07-07 19:58:21,158 - INFO - HTTP Request: POST https://claude.vocareum.com/v1/messages "HTTP/1.1 200 OK"
Let me check the database schema first:
[Using tool: read_query]
2026-07-07 19:58:21,159 - INFO - Attempting execution of tool 'read_query' (attempt 1/3)
2026-07-07 19:58:23,660 - INFO - HTTP Request: POST https://claude.vocareum.com/v1/messages "HTTP/1.1 200 OK"

[Using tool: read_query]
2026-07-07 19:58:23,662 - INFO - Attempting execution of tool 'read_query' (attempt 1/3)
2026-07-07 19:58:27,432 - INFO - HTTP Request: POST https://claude.vocareum.com/v1/messages "HTTP/1.1 200 OK"

[Using tool: read_query]
2026-07-07 19:58:27,434 - INFO - Attempting execution of tool 'read_query' (attempt 1/3)
2026-07-07 19:58:30,831 - INFO - HTTP Request: POST https://claude.vocareum.com/v1/messages "HTTP/1.1 200 OK"
Let me check what data is in the database:
[Using tool: read_query]
2026-07-07 19:58:30,831 - INFO - Attempting execution of tool 'read_query' (attempt 1/3)
2026-07-07 19:58:34,063 - INFO - HTTP Request: POST https://claude.vocareum.com/v1/messages "HTTP/1.1 200 OK"

[Using tool: read_query]
2026-07-07 19:58:34,065 - INFO - Attempting execution of tool 'read_query' (attempt 1/3)
2026-07-07 19:58:44,640 - INFO - HTTP Request: POST https://claude.vocareum.com/v1/messages "HTTP/1.1 200 OK"
Perfect! Here's a detailed comparison of **CloudRift AI** vs **DeepInfra** for DeepSeek V3:

## DeepSeek V3 Pricing Comparison

| Provider | Input Price (per 1M tokens) | Output Price (per 1M tokens) | Total Cost* |
|----------|-----------------------------|------------------------------|-------------|
| **CloudRift AI** | $0.27 | $1.10 | $1.37 |
| **DeepInfra** | $0.32 | $0.89 | $1.21 |
| **Difference** | +$0.05 (19% higher) | -$0.21 (19% lower) | -$0.16 (12% lower) |

*Total cost for 1M input + 1M output tokens

### Key Insights:

1. **Input Tokens**: CloudRift AI is **cheaper** at $0.27 vs DeepInfra's $0.32 (saving $0.05 or ~19%)

2. **Output Tokens**: DeepInfra is **significantly cheaper** at $0.89 vs CloudRift's $1.10 (saving $0.21 or ~19%)

3. **Overall Value**: DeepInfra offers better overall value if you have typical or output-heavy workloads

### Which to Choose?

- **Choose CloudRift AI** if: Your workload is heavily input-dominated (processing large documents with minimal response
s)
- **Choose DeepInfra** if: Your workload generates substantial output or has a balanced input/output ratio (most common
use case)

For a typical 50/50 input/output split (500K input, 500K output tokens):
- **CloudRift**: $0.135 + $0.55 = **$0.685**
- **DeepInfra**: $0.16 + $0.445 = **$0.605**
- **DeepInfra saves ~12%** in this scenario2026-07-07 19:58:51,053 - INFO - HTTP Request: POST https://claude.vocareum.c
om/v1/messages "HTTP/1.1 200 OK"
2026-07-07 19:58:51,056 - INFO - Stored 0 pricing plans

=== SHOWING STORED DATA ===
2026-07-07 19:58:51,056 - INFO - Attempting execution of tool 'read_query' (attempt 1/3)
```